



CouchDB – overview

Sebastian Rothbucher @sebredhh
with lots of help from @awenkhh


What is it - @all?

- Cluster of unreliable commodity hardware
- Map-Reduce & Document-oriented Database (no-SQL, no-schema)
- „Brainchild“ of a former senior IBM Lotus Domino Developer
- Web to the core (HTTP, JSON, JavaScript)
- Document = versioned JSON = unit of work = many roles
- Web-interface (and curl)
- Stable and scalable (Erlang-based foundation)

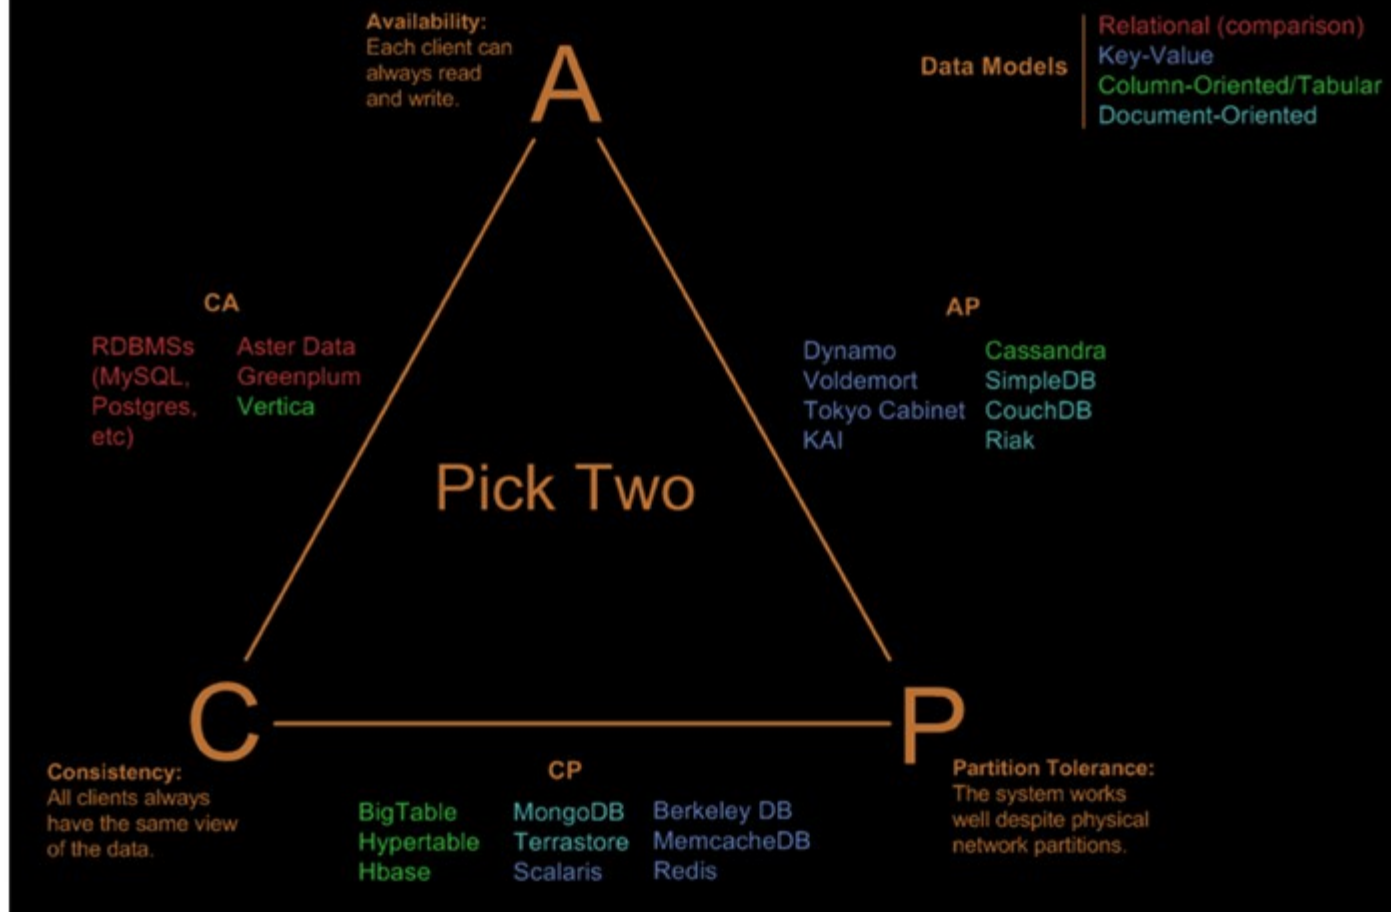
What you'll hear

- NoSQL database overview
- What is CouchDB?
- Querying CouchDB with http
- Fauxton
- Retrieving data with Map/Reduce
- MVCC
- `_changes` *magic*

How does it differ?



Visual Guide to NoSQL Systems



Web, web, web



```
curl -X PUT -d ...
```

Create DB & docs



Add attachment & serve infos



Futon > Fauxton & Docs



Map-Reduce („closures“)

Map (produce index)

- „Explode“ document base
- No-schema – hence check docs
- Produce keys/values (0-n)
- key, docid, value
- Sorted (B+Tree)
- Highly efficient & (somewhat) parallel
- (map-only possible)

Default: `_all_docs` (key = docid)

➔ Analyzed set is not base set
(query on index ONLY != SQL: on table)

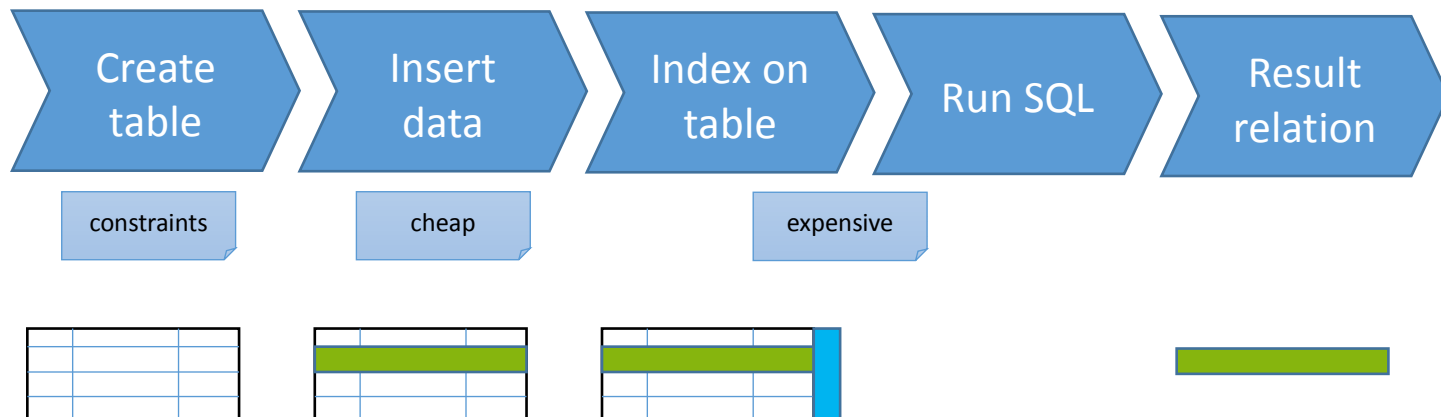
Reduce (aggregate index)

- Collated set (one key)
- Produce one output from several
- „Collapse“ Info
- Sum, Count, Stats
- (Complex (Correlation, max. one „join“, ...))

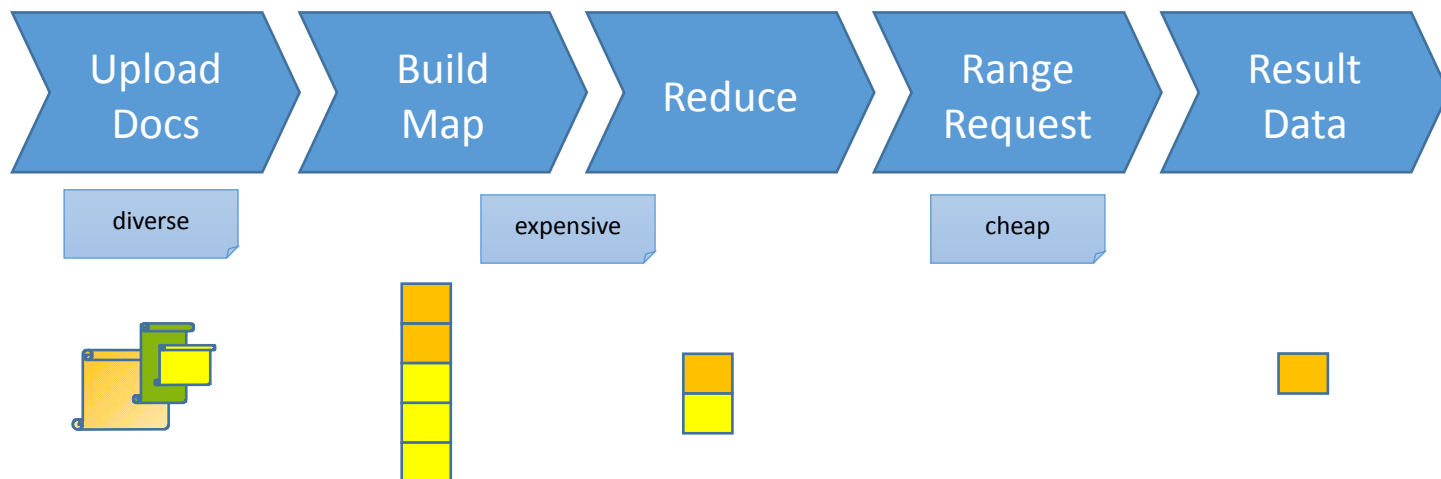
➔ Huge indices on one key

How does it fit together?

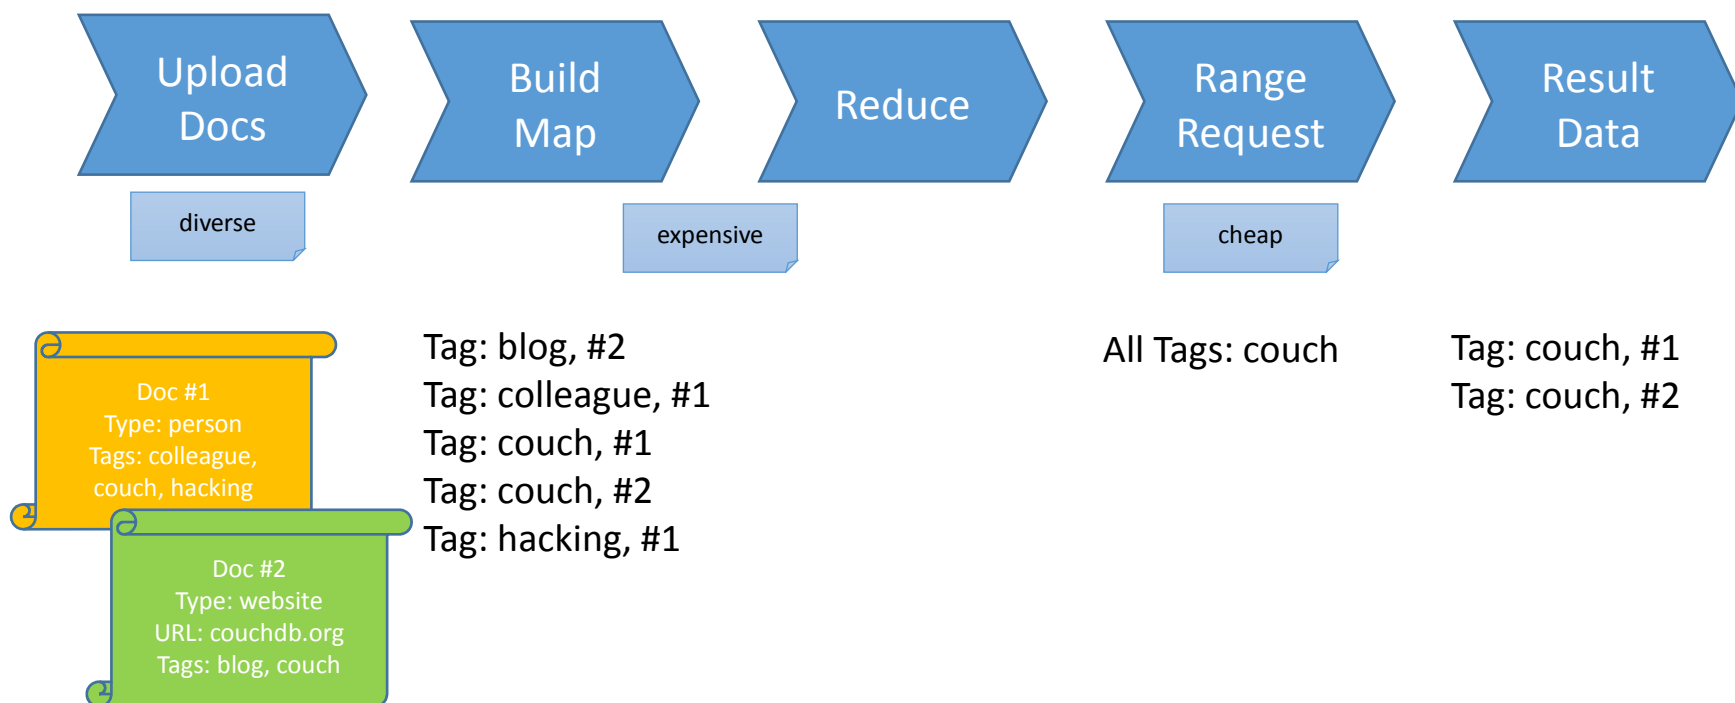
SQL
(relational)



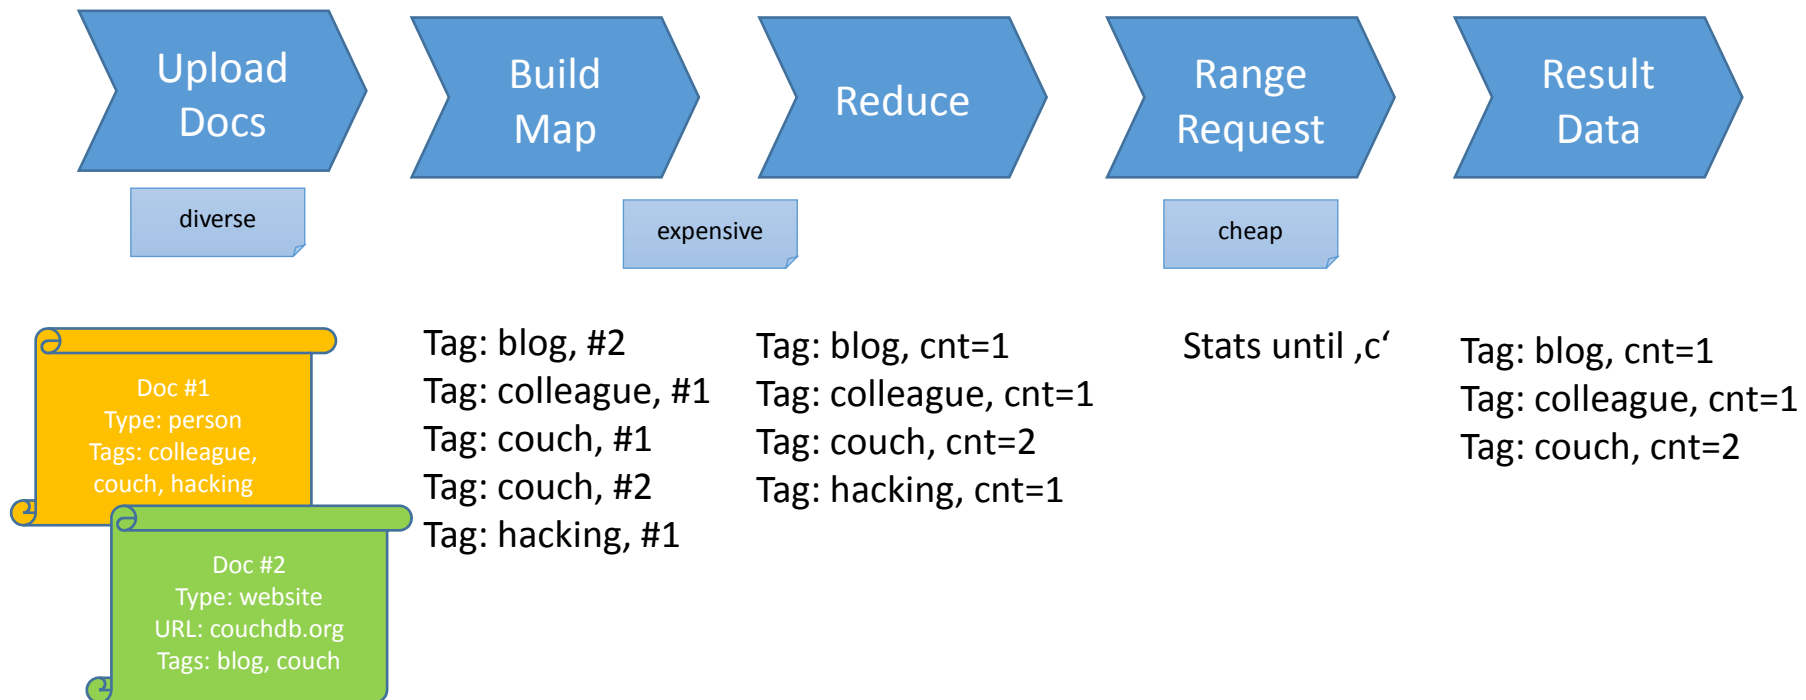
Couch
(M/R)



How does it fit together? (example)



How does it fit together? (example cont.)



Creating a view



Querying a view & paging






More curl magic (?)




Advanced: view collation (?)




MVCC




1. Update V=1... (OK)




MVCC (cont.)




 != versioning!


Replication



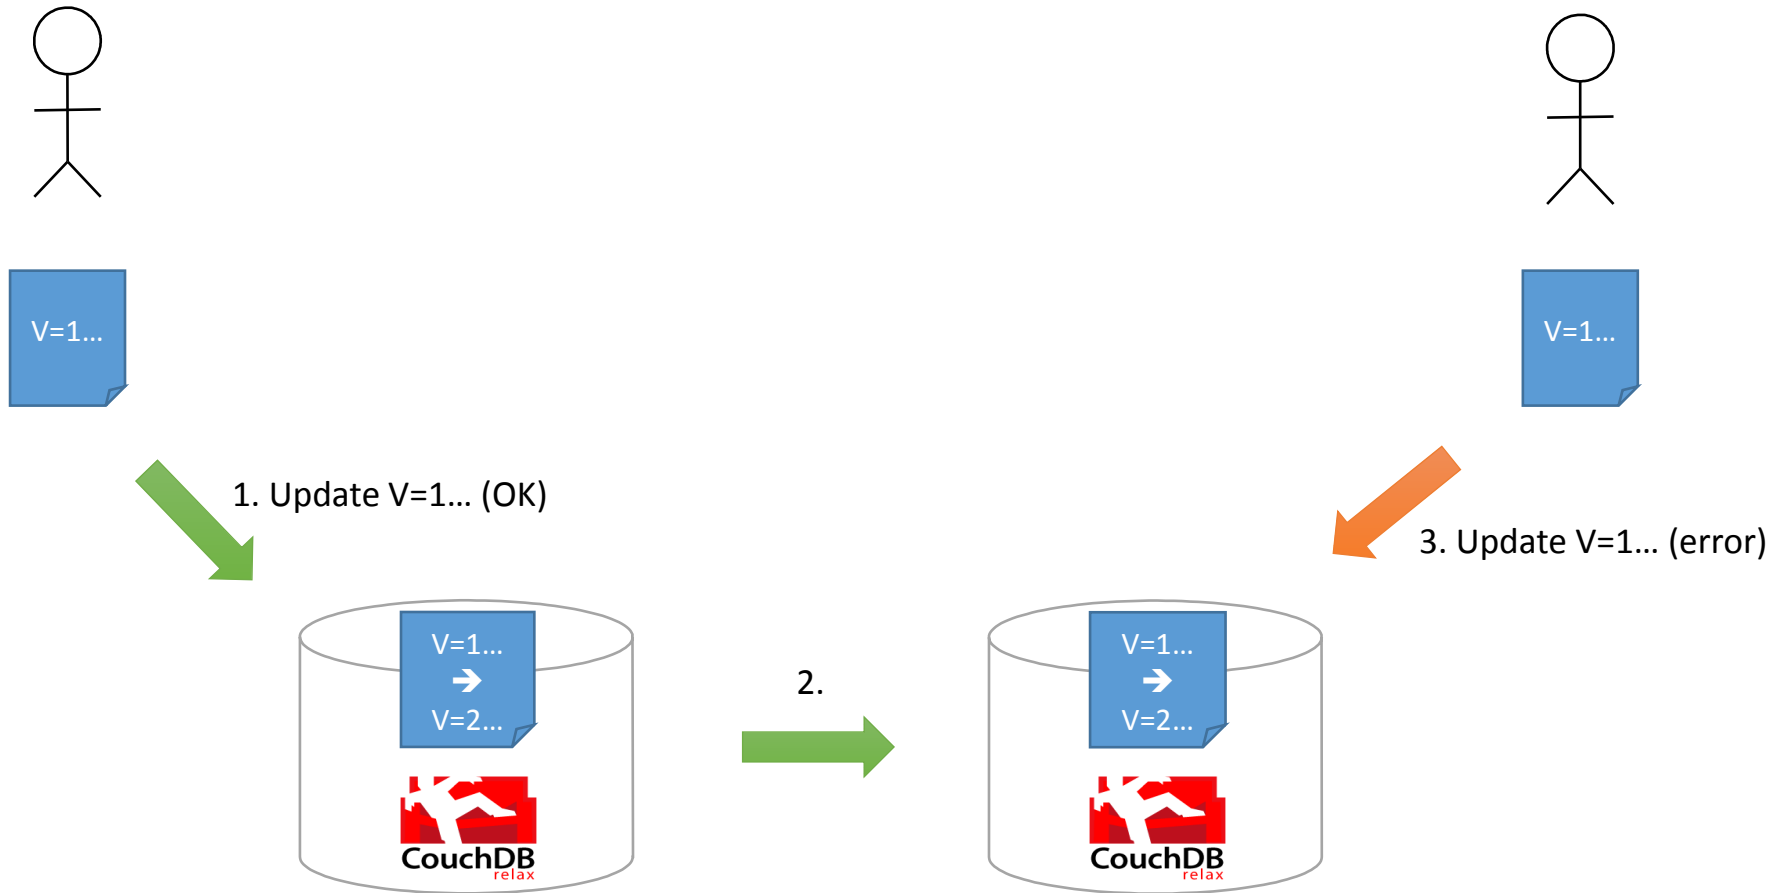
1. Update V=1... (OK)



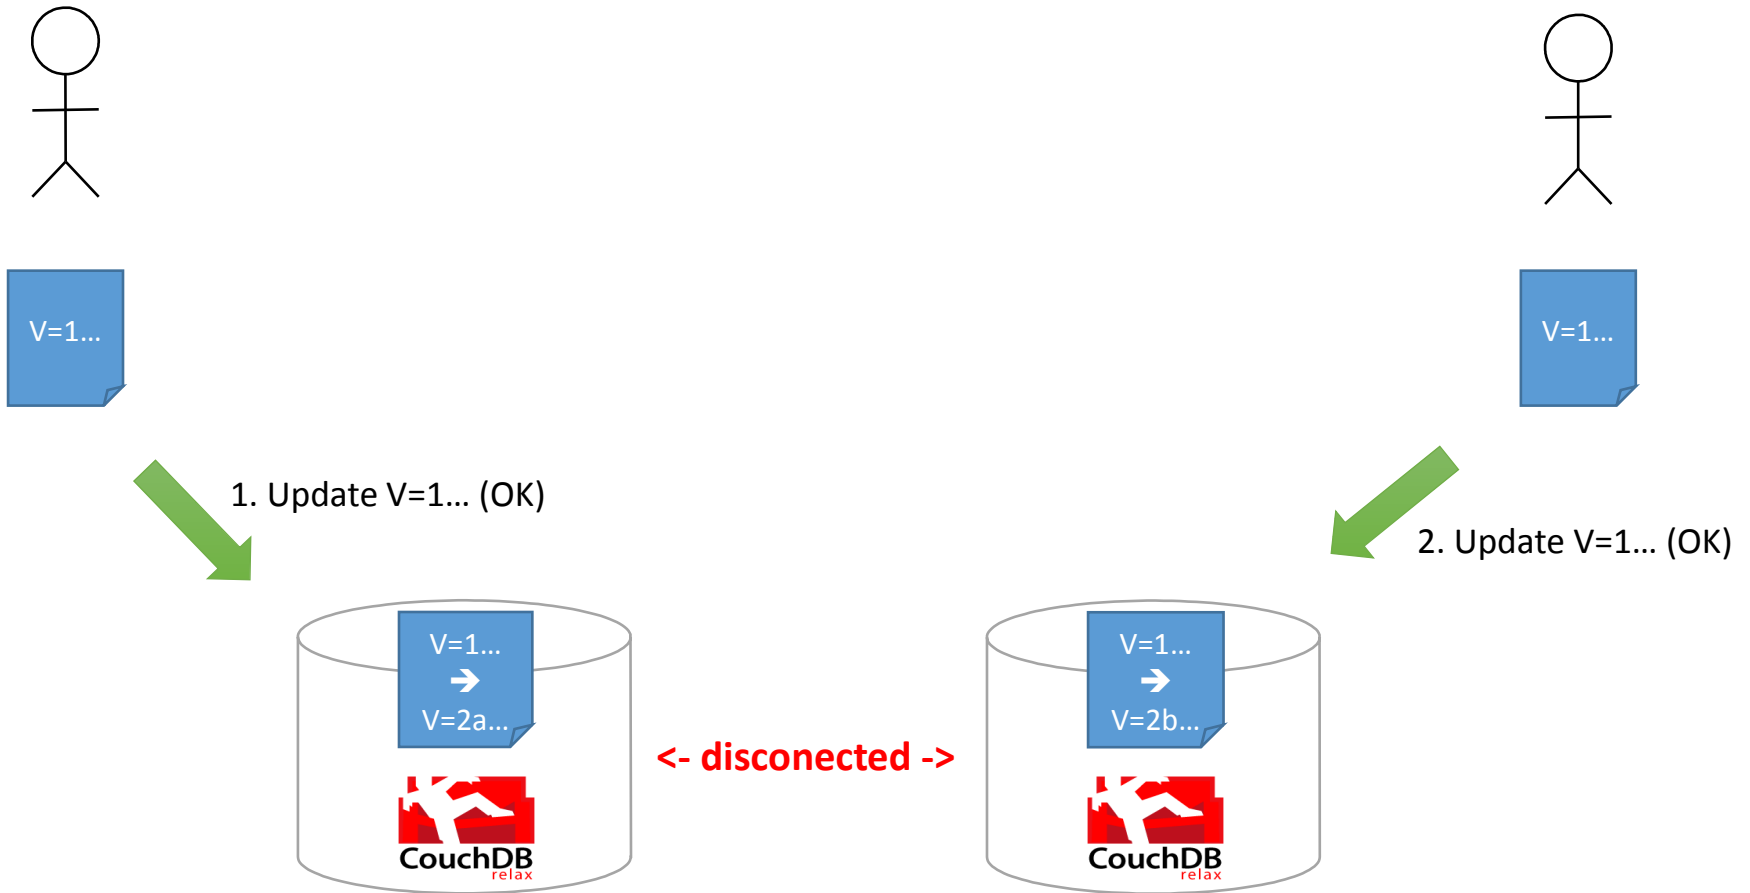
2.



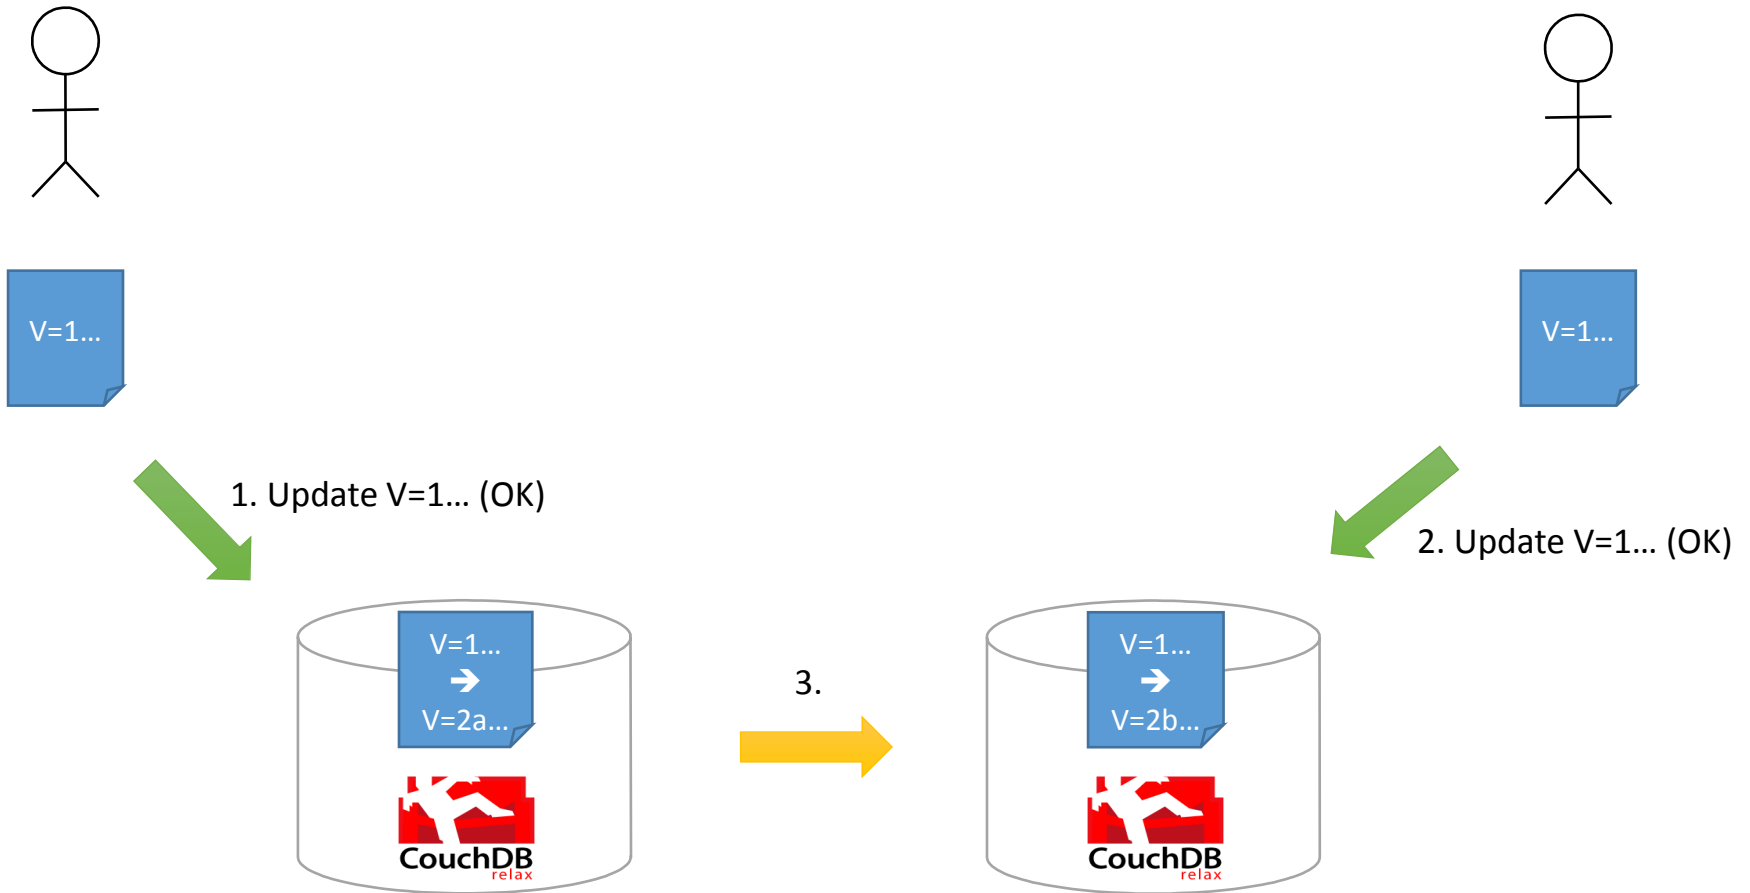
Replication (cont.)



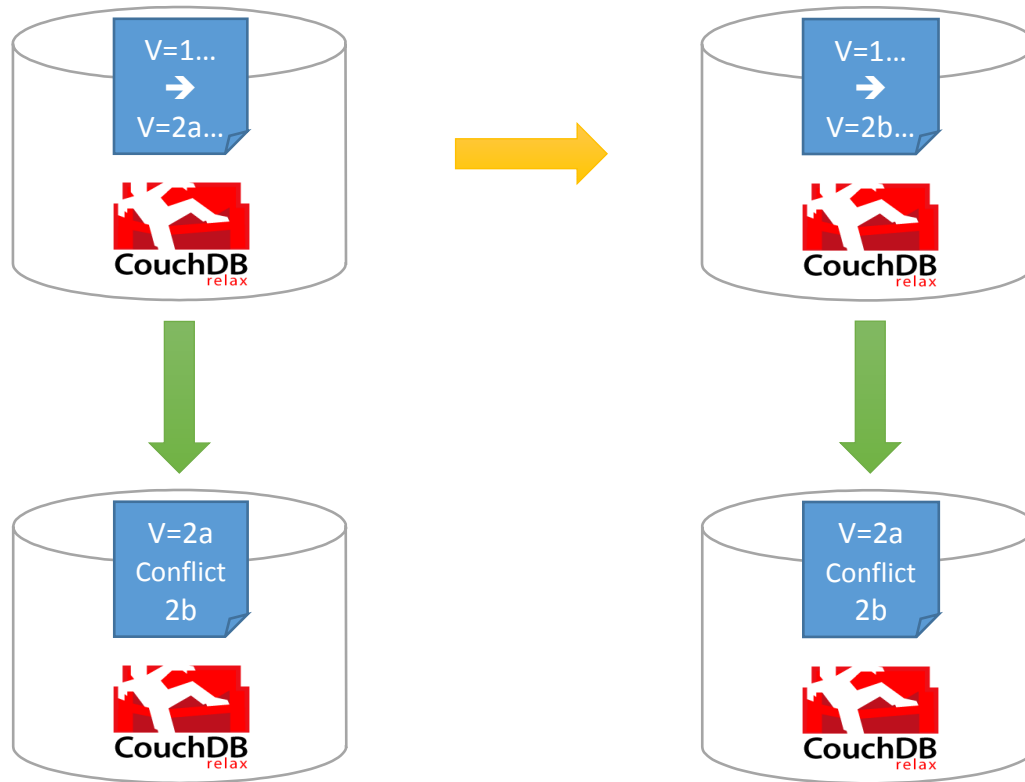
Replication (cont.)



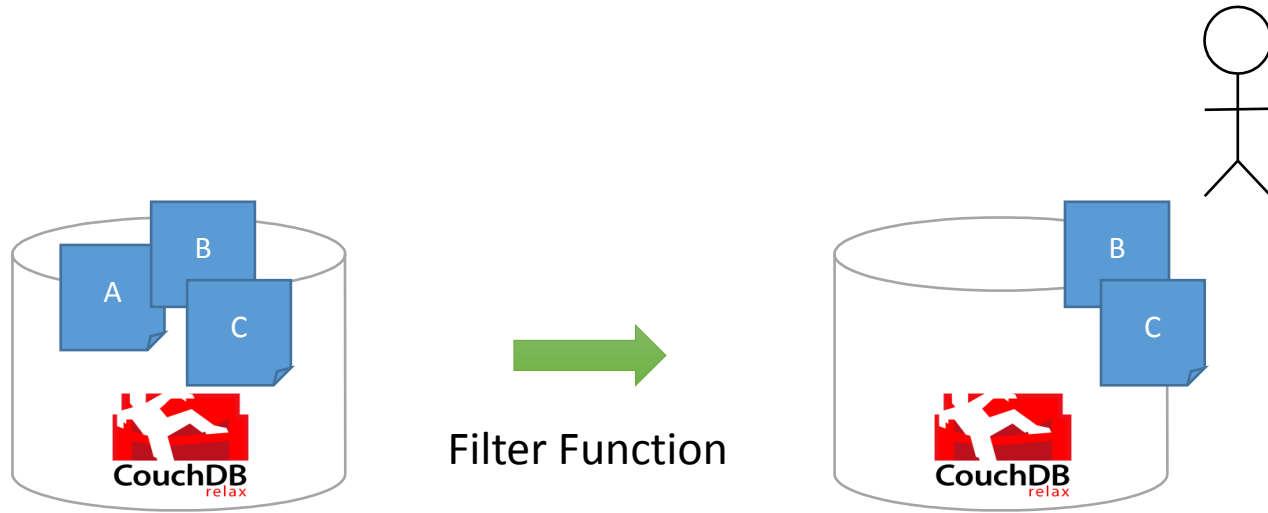
Replication (cont.)



Replication (cont.) – Conflict resolution

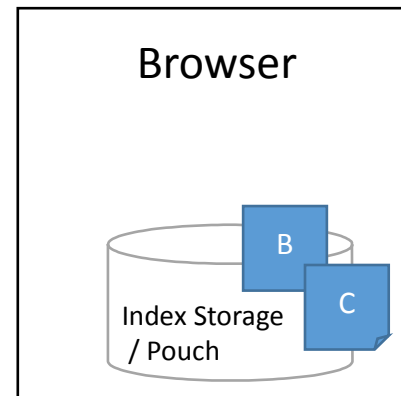




Replication - selective




➔ Couch per user


Replication PouchDB



Replication PouchDB (cont.)



<- disconnected ->



Doing some replication



The best is yet to come...



`_changes`

Links



- CouchDB Website: <http://couchdb.apache.org/>
- CouchDB day talks:
<http://blog.couchdb.org/2015/02/24/couchdb-day-2015-recap/>
- CouchDB Manual: <http://docs.couchdb.org/>
- Joan's talk on MVCC et al: <http://atypical.net/archive/2014/02/04/my-couchdb-conf-talk>
- eHealth Africa – collected best practices:
<http://docs.ehealthafrica.org/couchdb-best-practices/>
- Jan's intro into CouchDB:
<http://www.infoq.com/articles/apache-couchdb-the-definitive-introduction>
- Cloudant Docs: <https://docs.cloudant.com/>

Summary & Wrap-Up

```
http://  
    function map(doc)  
  
emit(doc.tags[i], null)  
  
?startkey=couch&limit=10  
  
<script src="angular.js"  
  
    _changes?feed=longpoll  
  
{ "couchdb": "Welcome"
```

JUST DO IT